
VisConf: Quality-Aware Consensus for Best-of-N in Vision-Language Model Reasoning

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Best-of-N methods improve text-only LLM reasoning by sampling multiple trajec-
2 tories before selecting an answer. For Vision-Language Models (VLMs), answer
3 frequency and model confidence alone can be insufficient, as sampled generations
4 also vary in their engagement with the visual input. We find that distributional
5 confidence and visual attention are complementary indicators of rollout quality.
6 Each is limited on its own, while trajectories that show strong signals in both are
7 more likely to be correct. Motivated by this, we introduce VisConf, a rollout-quality
8 score that multiplicatively combines visual engagement with candidate-calibrated
9 Self-Certainty, penalizing rollouts that are weak in either signal. We incorporate
10 Rank-Weighted Consensus (RWC), a tuning-free aggregation rule that converts
11 within-question VisConf ranks into answer-level vote weights. Across three VLMs
12 on MathVista, MMMU-Pro, and MMStar, VisConf with RWC consistently out-
13 performs strong BoN baselines. Our analysis shows particularly large gains when
14 correct trajectories are present but underrepresented in the candidate pool. [Code](#).

15 1 Introduction

16 A widely used test-time scaling method in Large Language Model (LLM) reasoning is Best-of-N
17 (BoN) selection, where multiple candidate trajectories are sampled for each query [Snell et al.,
18 2024, Brown et al., 2024]. Generating multiple rollouts using stochastic decoding provides more
19 opportunities to produce correct reasoning paths, but also introduces a selection problem since
20 repeated sampling can produce redundant reasoning patterns rather than genuinely diverse solutions
21 [Bilal et al., 2026]. Vision-Language Models (VLMs) similarly benefit from rollout scaling [Sammani
22 et al., 2026, Tong et al., 2026], but multimodal reasoning introduces another dimension of candidate
23 variation as trajectories can differ in how strongly they engage with the visual input.

24 Two natural signals for assessing VLM rollout quality are confidence and visual engagement. Output-
25 probability-based confidence captures how decisively the model supports its generated trajectory, and
26 has been widely used as a proxy for reliability [Geng et al., 2024, Kang et al., 2025]. However, high
27 confidence does not necessarily imply correct visual reasoning [Cho et al., 2026]. Visual attention can
28 indicate greater reliance on the image, but may also reflect repeated inspection caused by uncertainty
29 or misinterpretation. Our empirical analysis shows that these signals are weakly informative but
30 provide complementary information at the trajectory level: rollouts exhibiting both high confidence
31 and strong visual engagement are substantially more reliable than those strong in only one dimension.
32 We further find that their useful interaction is primarily trajectory-level rather than token-aligned.
33 Motivated by these findings, we introduce VisConf, which independently aggregates confidence and
34 visual engagement over a rollout before combining them into a single quality score.

35 For verifiable tasks, BoN exposes two sources of evidence: the quality of individual rollouts and
36 agreement across the candidate set. Self-Consistency exploits the latter by favoring answers supported

by multiple independently sampled trajectories, under the intuition that diverse reasoning paths converging to the same answer provide stronger evidence of correctness [Wang et al., 2023]. However, standard Self-Consistency treats all rollouts equally and fails when redundant incorrect generations form a majority group [Yin et al., 2024, Taubenfeld et al., 2025]. In contrast, conventional BoN selection relies only on the single highest-scoring rollout, discarding corroborating evidence from other trajectories that share the same outcome. We therefore use Rank-Weighted Consensus (RWC), a Borda-inspired hyperparameter-free aggregation scheme, to combine trajectory quality with answer-level consensus.

Our contributions are as follows: (1) We empirically characterize confidence and visual engagement as complementary trajectory-level signals of VLM rollout quality, showing that their useful interaction is primarily trajectory-level rather than token-aligned. (2) We introduce a VLM BoN framework featuring VisConf, which combines candidate-calibrated Self-Certainty with visual engagement, and RWC, which incorporates trajectory quality into answer-level consensus. (3) We demonstrate consistent improvements across three models, three benchmarks, and three sampling budgets, including average accuracy gains of 0.98 points over Self-Consistency and a statistically significant 0.49 points over the strongest baseline.

2 Related Work

Test-Time Scaling and Best-of-N Test-time scaling improves performance by allocating additional inference compute without modifying model parameters [Snell et al., 2024, Brown et al., 2024]. Self-Consistency aggregates candidates through answer frequency, but provides no explicit signal of rollout quality [Wang et al., 2023]. Alternatively, verifier- and reward-model-based methods rank candidates using learned correctness features, at the cost of training or invoking an additional model [Cobbe et al., 2021, Lightman et al., 2023]. Recent work in VLMs finds modest gains from single-model scaling while evaluating lexical/length-based heuristics and majority voting, leaving stronger model-internal rollout-quality signals unexplored [Tong et al., 2026].

Confidence-Based Selection Model confidence provides a verifier-free quality signal for sampled trajectories [Geng et al., 2024, Kang et al., 2025]. Likelihood-based measures such as sequence log-probability use only the probabilities assigned to the sampled tokens, ignoring the remainder of the predictive distribution [Kang et al., 2025, Kim and Kim, 2026]. Self-Certainty instead uses the full predictive distribution through Kullback-Leibler (KL) divergence from a uniform reference distribution [Kang et al., 2025]. However, its Borda-based aggregation introduces a tunable exponent.

Visual Grounding in VLMs VLM responses can be weakly image-conditioned, exhibiting hallucination and language-prior bias, motivating methods that promote visual faithfulness [Li et al., 2023, Leng et al., 2024]. VGPO identifies temporal visual forgetting, showing that sustained visual attention can be indicative of correctness [Wang et al., 2026]. Most closely related, LookBack combines token likelihood with attention-based visual reference to select the highest-scoring rollout [Cho et al., 2026]. In contrast, VisConf calibrates distributional confidence within the candidate set, aggregates confidence and visual engagement separately over each rollout before fusion, and uses RWC to retain answer-level consensus.

3 What Signals VLM Rollout Quality?

BoN succeeds only when the selection rule can distinguish reliable reasoning from plausible errors. During autoregressive generation, the transformer architecture exposes two distinct sources of information at each decoding step: its predictive distribution over the vocabulary, from which we derive *confidence*, and its internal attention to the visual input, which we term *visual engagement*.

3.1 Standalone Quality Signals

Confidence via Self-Certainty For a generation trajectory i at step t , let $p_{i,t}$ denote the next-token distribution obtained by applying softmax directly to the unscaled model logits, and U the uniform distribution over the vocabulary. The token-level confidence signal is

$$k_{i,t} = D_{\text{KL}}(U \parallel p_{i,t})$$




Figure 1: Analysis on Qwen2.5-VL-3B using MathVista, MMMU-Pro, and MMStar. (a) Accuracy as a function of confidence and visual engagement. (b) Discriminative strength measured by the accuracy gap between the top and bottom score quartiles. Combining confidence and visual engagement is more informative, but how the signals are combined also matters.

85 where larger values indicate more concentrated predictive distributions [Kang et al., 2025].

86 **Visual Engagement via Visual Attention** We use visual attention as a token-level measure of
 87 visual engagement, following LookBack [Cho et al., 2026]. Let $\alpha_{i,t,j}^{(\ell,h)}$ be the post-softmax attention
 88 from the position predicting token t to context position j in layer ℓ and head h . We average attention
 89 across all layers and heads and sum within each token group S :

$$\bar{\alpha}_{i,t,j} = \frac{1}{\sum_{\ell} H_{\ell}} \sum_{\ell} \sum_{h=1}^{H_{\ell}} \alpha_{i,t,j}^{(\ell,h)}, \quad A_{i,t}^S = \sum_{j \in S} \bar{\alpha}_{i,t,j}.$$

90 For image ($\mathcal{I}$), prompt-text ($\mathcal{P}$), and all previously generated-token ($\mathcal{G}_i$) positions, excluding padding
 91 and non-image special tokens, we define token-level visual engagement as

$$v_{i,t} = \frac{A_{i,t}^{\mathcal{I}}}{A_{i,t}^{\mathcal{I}} + A_{i,t}^{\mathcal{P}} + A_{i,t}^{\mathcal{G}_i}}.$$

92 Larger $v_{i,t}$ indicates a greater share of attention allocated to the image relative to textual context.

93 We then average the token-level signals for a generation i to obtain raw confidence and visual
 94 engagement scores:

$$K_i = \frac{1}{T_i} \sum_{t=1}^{T_i} k_{i,t}, \quad V_i = \frac{1}{T_i} \sum_{t=1}^{T_i} v_{i,t}.$$

95 Top-1 BoN selection chooses rollout $\hat{i}$ among N rollouts using $\hat{i} = \arg \max_{i \in [1,N]} Q_i$, where
 96 $Q_i = K_i$ for confidence-based selection and $Q_i = V_i$ for visual engagement-based selection.

97 3.2 Each Signal is Informative But Weak

98 We first examine how strongly confidence and visual engagement individually predict correctness. We
 99 measure how well trajectory-level scores separate correct from incorrect rollouts using the accuracy
 100 gap between the top and bottom score quartiles. A useful quality signal should concentrate more
 101 correct rollouts in its highest-scoring quartile. We additionally report macro-averaged prompt-level
 102 AUC in the appendix (Table 6).

103 Mean confidence, K_i , produces an average top-bottom accuracy gap of 3.5 percentage points (Fig-
 104 ure 1b). High confidence may reflect a coherent and internally consistent reasoning path, but VLMs
 105 can also be confidently wrong when language priors dominate or an early visual misinterpretation
 106 propagates through otherwise fluent reasoning. Mean visual engagement, V_i , yields a slightly larger
 107 gap of 4.8 points. Sustained image attention may indicate continued use of visual evidence, but can




Figure 2: Overview of the VisConf and RWC framework. Given multiple sampled VLM rollouts, VisConf aggregates calibrated confidence and visual engagement independently over each trajectory, and combines them into a rollout-quality score, rewarding trajectories that are both confident and visually engaged. RWC uses VisConf scores to effectively select high-quality answer groups.

also reflect uncertainty, repeated reinspection, or attention to irrelevant content. Conversely, some correct rollouts may also extract necessary visual information during early reasoning and therefore not require prolonged visual engagement.

3.3 Confidence and Vision are Complementary at the Trajectory Level

Joint Signal Structure The two signals become notably informative when considered jointly. We stratify rollouts by whether their within-prompt raw-confidence and visual engagement scores fall above or below the median. Rollouts high in both signals achieve 49.5% accuracy, compared with 48.3% when only visual engagement is high, 47.9% when only confidence is high, and 43.4% when both are low (Figure 1a). Holding either signal fixed, higher values of the other are also associated with greater correctness enrichment (Appendix Figure 6). Together, these results indicate that confidence and visual engagement provide complementary rather than redundant evidence of rollout quality.

Trajectory-Level vs. Token-Level Fusion The joint analysis motivates combining confidence and visual engagement, but does not specify where fusion should occur. Token-level fusion averages the stepwise products, $Q_i^{\text{tok}} = \frac{1}{T_i} \sum_t v_{i,t} k_{i,t}$, coupling the two signals at the same decoding steps. Trajectory-level fusion instead multiplies their rollout means, $Q_i^{\text{traj}} = V_i K_i$, allowing confidence and visual engagement to contribute at different stages of generation. Both improve separation over either singleton signal, but trajectory-level fusion produces a larger top-bottom quartile accuracy gap than token-level fusion (6.0 vs. 5.4 percentage points; Figure 1b).

The two constructions differ only through within-rollout temporal covariance: trajectory-level fusion retains the marginal strength of each signal without requiring their tokenwise peaks to coincide. Consistent with this decomposition, randomizing token correspondence matches the AUC of trajectory-level fusion, whereas aligned token fusion is slightly weaker (Appendix D.2). This suggests that the useful interaction between confidence and visual engagement lies primarily in their trajectory-level magnitudes.

4 VisConf: Trajectory-Level Visual-Confidence Scoring

Shaped by our analysis in Section 3, we introduce VisConf, which first calibrates confidence relative to the candidate set and then combines rollout-level confidence and visual engagement (Figure 2).

136 4.1 Candidate-Relative Confidence Calibration

137 To integrate Self-Certainty into a multi-signal score, we calibrate confidence relative to the pooled
138 token-level values across all candidate rollouts G for a query.

$$c_{i,t} = \text{clip}_{[0,1]} \left(\frac{k_{i,t} - Q_{.05}(G)}{\max(Q_{.95}(G) - Q_{.05}(G), \epsilon)} \right),$$

139 where $Q_p(G)$ is the p -quantile of pooled token KL values in G , and $\epsilon > 0$ ensures numerical
140 stability. This scales confidence to $[0, 1]$ while bounding the influence of extreme KL values. Using
141 shared percentiles across G therefore situates each rollout in the context of the candidate set, making
142 subsequent trajectory-level scores relative rather than absolute.

143 4.2 Trajectory-Level Visual-Confidence Fusion

144 VisConf aggregates visual engagement and calibrated confidence independently over a complete
145 rollout by taking the average across the full generation, resulting in V_i and C_i , respectively. We then
146 define the VisConf quality score as:

$$Z_i = V_i C_i.$$

147 VisConf favors trajectories that exhibit both high calibrated confidence and high visual engagement,
148 without explicitly requiring that the two signals coincide at the same decoding steps.

149 5 From Rollout Quality to Answer Consensus

150 VisConf estimates the relative quality of each rollout, but selecting only the highest-scoring candidate
151 treats this noisy score as a verifier and ignores corroborating evidence from other trajectories.
152 The value of retaining both sources of evidence is reflected in their joint association with rollout
153 correctness (Appendix Figure 8). Inspired by Borda voting [Kang et al., 2025], we present *Rank-
154 Weighted Consensus* (RWC) to consolidate consensus with quality, as illustrated in Figure 2(c).

155 5.1 Rank-Based Quality Weighting

156 Let r_i denote the rank of a quality score Z_i among N rollouts, with $r_i = 1$ being the highest-ranked
157 candidate. We convert each rank into a normalized quality weight:

$$w_i = \frac{N - r_i}{N - 1},$$

158 with exact score ties receiving their average rank. The highest-ranked rollout receives weight 1 and
159 the lowest receives 0. Lower-ranked rollouts contribute proportionally less. Because these weights
160 depend only on ordering, they are invariant to strictly increasing transformations of the VisConf score.
161 Unlike Self-Certainty’s parameterized Borda weights $b_p(r_i) = (N - r_i + 1)^p$, RWC uses a fixed
162 linear schedule that downweights weak evidence without requiring p to be tuned [Kang et al., 2025].

163 5.2 Quality-Weighted Answer Consensus

164 Let a_i denote the normalized answer extracted from rollout i . We sum the quality weights of all
165 trajectories supporting each answer and select the answer with the highest quality-informed support:

$$\text{Vote}(a) = \sum_{i=1}^N \mathbf{1}[a_i = a] w_i, \quad \hat{a} = \arg \max_a \text{Vote}(a).$$

166 Thus, repeated answers accumulate support, but higher-quality trajectories exert more influence on
167 the final prediction. A minority answer can therefore prevail when its quality compensates for its
168 lower frequency.

169 6 Experimental Setup

170 We evaluate response-selection methods in VLMs under a shared evaluation protocol.

Table 1: Averaged accuracy (%) across MathVista, MMMU-Pro, and MMStar at varying rollout budgets N . Per-dataset results are presented in Appendix B.1. VisConf + RWC achieves the highest average for every model and sampling budget.

Method	Qwen2.5-VL-3B			InternVL3-8B			Gemma4-E2B			Avg.
	$N = 8$	$N = 16$	$N = 32$	$N = 8$	$N = 16$	$N = 32$	$N = 8$	$N = 16$	$N = 32$	
FirstAns	48.51	48.51	48.51	56.55	56.55	56.55	47.34	47.34	47.34	50.80
Self-Consistency	53.70	54.75	55.07	60.96	61.71	61.80	49.85	50.23	50.35	55.38
Self-Certainty	49.78	49.95	49.94	58.83	59.03	59.07	49.43	49.62	49.81	52.83
+ Borda ($p = 0.3$)	53.83	54.87	55.14	61.21	61.97	62.25	50.12	50.50	50.71	55.62
+ Borda ($p = 1.2$)	53.53	54.72	55.33	61.21	62.03	62.78	50.35	50.82	50.91	55.74
LookBack	50.09	50.34	50.57	58.22	58.11	57.84	49.16	49.74	50.27	52.70
+ RWC	53.68	54.93	55.72	61.27	62.16	62.65	50.30	50.86	51.19	55.86
VisConf	51.04	51.16	51.17	59.46	59.81	60.27	50.07	50.49	50.58	53.78
+ RWC	54.09	55.32	56.29	61.69	62.53	63.27	50.89	51.44	51.68	56.36

Models and Benchmarks: We evaluate Qwen2.5-VL-3B [Bai et al., 2025], InternVL3-8B [Zhu et al., 2025], and Gemma4-E2B [Gemma Team et al., 2026] on three multimodal reasoning benchmarks: MathVista [Lu et al., 2024], MMMU-Pro [Yue et al., 2025], and MMStar [Chen et al., 2024], using the MathVista *testmini* split and the full evaluation sets of MMMU-Pro and MMStar.

Candidate Generation and Evaluation: For each question, we generate 32 chain-of-thought rollouts [Wei et al., 2022, Zhang et al., 2024] using a fixed stochastic decoding configuration across models. We use temperature = 0.7, top- $p = 0.8$, top- $k = 20$, and a maximum of 2048 generated tokens. We report selection accuracy for $N \in \{8, 16, 32\}$. For $N = 8$ and $N = 16$, results are averaged over 100 fixed subsets sampled without replacement from the 32-rollout pool using fixed seeds. $N = 32$ uses the complete pool once. Accuracy is averaged within each question before dataset-level aggregation.

Baselines: We compare against *FirstAns*, which is the expected accuracy of one randomly sampled rollout, and *Self-Consistency*, which selects the most frequent answer. We evaluate *Self-Certainty* using top-1 BoN selection as well as Borda-based aggregation with $p \in \{0.3, 0.7, 1.2, 2.0\}$ ¹ [Kang et al., 2025]. For multimodal scoring, we evaluate *LookBack* [Cho et al., 2026] using its original top-1 selection and additionally pair its scores with RWC, allowing comparison under the same answer-aggregation rule. Finally, we evaluate VisConf with top-1 selection and RWC, recomputing confidence normalization and ranks within each sampled candidate subset.

7 Results

7.1 Overall Performance and Scaling

Table 1 reports accuracy averaged across the three benchmarks. VisConf + RWC achieves the highest accuracy for every model and sampling budget, with an overall average of 56.36%. This exceeds LookBack + RWC, the strongest baseline overall, by 0.5 points; the best Self-Certainty + Borda configuration by 0.62 points; and Self-Consistency by 0.98 points. Across the per-dataset results (Appendix Table 4), VisConf + RWC outperforms every baseline in 24 of 27 model-dataset-budget settings and trails by at most 0.27 points in the remaining three. Improvement against the strongest baseline remains statistically significant (paired 95% CI: [0.32, 0.66]) and positive at every sampling budget (Appendix B.2).

Comparing top-1 selection without aggregation, VisConf achieves an overall average of 53.78%, 1.08 and 0.95 points higher than LookBack (52.70%) and Self-Certainty (52.83%), respectively. VisConf’s advantage over LookBack suggests that independently high confidence and visual engagement across decoding stages yields a stronger quality signal than requiring both signals to coincide token-wise. These top-1 results encourage the potential use of VisConf where aggregation schemes are not applicable. In our settings, however, omitting answer-group information is generally restrictive, as

¹Results for $p = 0.7$ and $p = 2.0$ are reported in Appendix B.1.




Figure 3: Effect of scaling N on accuracy across three benchmarks. (a) VisConf top-1 selection scales strongly but underperforms Self-Consistency. (b-c) VisConf + RWC consistently outperforms top-1 and Self-Consistency, with its advantage persisting or increasing with scaling.

top-1 VisConf selection performs below VisConf + RWC across all 27 settings. This validates the use of quality-informed aggregation in tasks where outcome-level rollout grouping is available.

VisConf + RWC provides additional benefits as N is scaled. From $N = 8$ to $N = 32$, VisConf + RWC improves from 54.09% to 56.29% on Qwen2.5-VL-3B, 61.69% to 63.27% on InternVL3-8B, and 50.89% to 51.68% on Gemma4-E2B (Figure 3). Our formulation at $N = 16$ outperforms Self-Consistency at $N = 32$ in all three models, indicating more effective use of the rollout budget.

7.2 Ablation Study

Score Ablation Table 2 shows the results of isolating each VisConf score component under top-1 selection on Qwen2.5-VL-3B. Visual engagement and confidence without fusion are relatively weak selectors, at 49.58% and 49.89%, while their raw trajectory-level product improves performance to 50.12%. Incorporating candidate-relative calibration with token-level fusion achieves 50.44%, lower than our VisConf trajectory-level formulation, which is 0.68 points higher at 51.12%. These results validate individual components of VisConf as meaningfully contributing to the overall performance. Our VisConf score can be generally written as

$$Z_{i,\alpha} = V_i^\alpha C_i^{1-\alpha},$$

where α controls the relative influence of visual engagement and confidence. Since $Z_{i,0.5} = \sqrt{V_i C_i}$, it induces exactly the same ranking as the product $V_i C_i$ used by VisConf. A post-hoc sensitivity sweep validates this balanced choice without per-dataset tuning (Appendix Figure 7).

Furthermore, to assess whether VisConf’s performance can be attributed to generation length bias, we include a MinLength baseline that selects the shortest rollout. Its substantially lower accuracy (49.20%) suggests that VisConf’s gains are not driven by a simple preference for shorter generations.

Aggregation Ablation Holding the quality score fixed and replacing RWC with raw-score weighted consensus WC (where $w_i = Z_i$) reduces accuracy from 55.23% to 55.02%: ranks preserve the candidate ordering while discarding unreliable score magnitudes.

Table 2: Ablation study on Qwen2.5-VL-3B. Average accuracy % across the three benchmarks for $N \in \{8, 16, 32\}$.

Method	Avg.
Score Ablation with BoN	
Visual Engagement: V	49.58
Self-Certainty: K	49.89
VisConf w/o KL Calibration: VK	50.12
Token Fusion: $\overline{v_{i,t} C_{i,t}}$	50.44
VisConf: VC	51.12
MinLength	49.20
Aggregation Ablation	
VisConf + WC	55.02
VisConf + RWC	55.23




Figure 4: Performance comparison against Self-Consistency on InternVL3-8B in cases where the candidate pool contains few correct rollouts. VisConf + RWC effectively prioritizes correct evidence without majority support.

7.3 Performance When Correct Evidence Is Sparse

Figure 4 examines how performance varies with empirical generation difficulty, measured by the prevalence of correct-answer trajectories in the $N = 32$ candidate pool,

$$g_q = \frac{n_{\text{correct}}}{N}.$$

VisConf + RWC increasingly outperforms Self-Consistency as correct evidence approaches, but remains below, the absolute majority threshold. Its advantage rises from 0.96 percentage points with 1-4 correct rollouts to 4.79 points with 5-8, 10.29 points with 9-12, and 12.18 points with 13-15. All 110 questions rescued relative to Self-Consistency occur within $0 < g_q < 0.5$, compared with 29 regressions, yielding 81 net corrections in this subset (Appendix Table 7). Thus, VisConf + RWC provides its largest gains when substantial correct evidence exists but remains outnumbered under unweighted consensus, using quality information to prioritize stronger evidence over frequency alone.

8 Discussion

Best-of-N is both a discovery and an aggregation problem. Increasing the rollout budget results in two competing outcomes: it increases the chance of discovering a correct reasoning path, but is likely to introduce additional incorrect candidates which may also be high-scoring. Our direct-selection results demonstrate that simply extending the candidate set does not guarantee better BoN selection, even when the oracle accuracy of the pool continues to increase (Appendix B.1). The stronger scaling of VisConf + RWC suggests that the value of additional test-time compute depends not only on generating better candidates, but also on preserving useful evidence distributed across the candidate set. We observe that this preservation can positively affect outcomes in cases where correct reasoning is less commonly encountered.

Relative quality is more reliable than absolute quality. Confidence and visual engagement are noisy proxies, and neither of them provides a reliable estimate that a rollout is correct. VisConf and RWC therefore use these signals comparatively while carefully avoiding spurious conclusions. Candidate-relative calibration contextualizes confidence within the current pool, multiplicative fusion promotes the occurrence of both signals, and RWC retains only ordinal rankings rather than interpreting raw magnitudes. This separation of roles is important. VisConf provides a general measure of which trajectories appear stronger, while RWC gauges how this imperfect ordering aligns with answer-level agreement. An extreme raw score’s domination is therefore subdued (Appendix E.3) and any single trajectory’s outcome is not treated as conclusive. Our work suggests that effective selection is possible without requiring a perfect verifier. Weak model-internal signals can become useful provided that the influence of noise is systematically reduced.

Limitations and Future Work. While we incorporate attention as a reference for the model engaging with input image(s), attention alone does not establish correct visual understanding or faithfully explain the model’s reasoning [Jain and Wallace, 2019, Wiegrefe and Pinter, 2019]. Obtaining visual-attention signals also requires eager attention extraction. In a controlled replay study on Qwen2.5-VL-3B, the median runtime overhead relative to FlashAttention-2 was negligible at $N = 8$ and $N = 16$ but increased to 18.4% at $N = 32$ (Appendix C). Future work could replace attention with a cheaper proxy or one that more causally explains whether visual evidence is reliably used. Although VisConf’s superior top-1 performance over other scoring baselines suggests that it may remain useful when answer-level grouping is unavailable, we do not evaluate this setting directly. More broadly, VisConf provides a training-free estimate of rollout quality that need not be confined to selection, and future work could study its use as a reward signal for policy optimization or in other paradigms where judging trajectory quality is useful.

9 Conclusion

We introduced VisConf, a simple, training-free response quality score using model-internal quality signals. Together with RWC, our framework combines rollout quality with answer-level agreement to improve BoN selection while requiring no tunable hyperparameters or external reward models. Our results demonstrate that imperfect model-internal signals can support effective test-time scaling when used comparatively and accumulated across agreeing trajectories rather than treated as standalone verifiers.

References

- Shuai Bai, Keqin Chen, Xuejing Liu, Jialin Wang, Wenbin Ge, Sibao Song, Kai Dang, Peng Wang, Shijie Wang, Jun Tang, Humen Zhong, Yuanzhi Zhu, Mingkun Yang, Zhaohai Li, Jianqiang Wan, Pengfei Wang, Wei Ding, Zheren Fu, Yiheng Xu, Jiabo Ye, Xi Zhang, Tianbao Xie, Zesen Cheng, Hang Zhang, Zhibo Yang, Haiyang Xu, and Junyang Lin. Qwen2.5-VL technical report, 2025. URL <https://arxiv.org/abs/2502.13923>.
- Ahsan Bilal, Muhammad Ahmed Mohsin, Muhammad Umer, Lena Trigg, Ali Subhan, Muhammad Ali, and Dean F. Hougen. Refining Over Resampling: Test-Time Self-Correction for LLM Reasoning, 2026. URL <https://arxiv.org/abs/2608.05643>.
- Bradley Brown, Jordan Juravsky, Ryan Ehrlich, Ronald Clark, Quoc V. Le, Christopher Ré, and Azalia Mirhoseini. Large Language Monkeys: Scaling Inference Compute with Repeated Sampling, 2024. URL <https://arxiv.org/abs/2407.21787>.
- Lin Chen, Jinsong Li, Xiaoyi Dong, Pan Zhang, Yuhang Zang, Zehui Chen, Haodong Duan, Jiaqi Wang, Yu Qiao, Dahua Lin, and Feng Zhao. Are We on the Right Way for Evaluating Large Vision-Language Models?, 2024. URL <https://arxiv.org/abs/2403.20330>.
- Beomsik Cho, Jinhyeong Kim, Dongseok Lee, and Jaehyung Kim. LookBack: Where and How to Score LVLM Responses via Visual Reference Usage, 2026. URL <https://arxiv.org/abs/2608.11847>.
- Karl Cobbe, Vineet Kosaraju, Mohammad Bavarian, Mark Chen, Heewoo Jun, Lukasz Kaiser, Matthias Plappert, Jerry Tworek, Jacob Hilton, Reiichiro Nakano, Christopher Hesse, and John Schulman. Training Verifiers to Solve Math Word Problems, 2021. URL <https://arxiv.org/abs/2110.14168>.
- Gemma Team, Sherif El Abd, Vaibhav Aggarwal, Robin Algayres, Alek Andreev, Olivier Bachem, Ian Ballantyne, Cormac Brick, Victor Cărbune, Michelle Casbon, Mayank Chaturvedi, Aditya Chawla, Victor Cotruta, Alice Coucke, Phil Culliton, Robert Dadashi, Lucas Dixon, Mohamed Elhawaty, Utku Evci, Clément Farabet, Johan Ferret, Filippo Galgani, Sertan Girgin, Jean-Bastien Grill, Maarten Grootendorst, Jiaxian Guo, Cassidy Hardin, Yanzhang He, Steven M. Hernandez, Omri Homburger, Léonard Hussenot, Juyeong Ji, Armand Joulin, Aishwarya Kamath, Parnian Kassraie, Olivier Lacombe, Preethi Lahoti, Gaël Liu, Gus Martins, Luciano Martins, Tatiana Matejovicova, Ramona Merhej, Nikola Momchev, Sneha Mondal, Ryan Mullins, Sindhu Raghuram Panyam, Shreya Pathak, Sarah Perrin, André Susano Pinto, Etienne Pot, Angéline Pouget, Alexandre Ramé, Sabela Ramos, Douglas Reid, David Rim, Morgane Rivi re, Karsten Roth, Louis Rouillard, Omar Sanseviero, Pier Giuseppe Sessa, Shane Settle, Danila Sinopalnikov, Sara Smoot, Piotr Stanczyk, Andreas Steiner, Lawrence Stewart, Ilya Tolstikhin, Michael Tschannen, Anton Tsitsulin, Nino Vieillard, Renjie Wu, Pingmei Xu, Haichuan Yang, Edouard Yvinec, Biao Zhang, Li Zhang, Joe Zou, Nicolas Aagnes, Abdelrahman Abdelhamed, Jakub Adamek, Shivani Agrawal, Shubham Agrawal, Ibrahim Alabdulmohsin, Jean Baptiste Alayrac, Uri Alon, Chandramouli Amarnath, Ankesh Anand, Chrysovalantis Anastasiou, Setareh Ariafar, Fran ois-Xavier Aubet, Kyriakos Axiotis, Federico Barbero, Joelle Barral, Alexei Bendebury, Urs Bergmann, Stanley Bileschi, Kat Black, Mathieu Blondel, Sebastian Borgeaud, Arthur Bra zinskas, Ryan Burnell, Robert Busa-Fekete, Mu Cai, Daniele Calandriello, Glenn Cameron, Charlotte Caucheteux, Rahma Chaabouni, Garima Chadha, Jetha Chan, Blake Jianhang Chen, Jesse Chen, Lin Chen, Xu Chen, Derek Cheng, Tzu hsiang Chien, Nikolai Chinaev, Yi Chou, Zhaohui Chu, Benjamin Coleman, Pooja Consul, Sam Conway-Rahman, Scott Crowell, Dylan Cutler, Vivek Dani, Samira Daruki, Anil Das, Daniel Deutsch, Nishanth Dikkala, Li Ding, Qiuhuan Ding, Shenil Dodhia, Konstantin Donhauser, Tulsee Doshi, Anca Dragan, Alex Druinsky, Sahil Dua, Zoltan Egyed, Danielle Eisenbud, Daniel Eppens, Cindy Fan, Bahare Fatemi, Yassir Fathullah, Vlad Feinberg, Milen Ferev, Sebastian Flennerhag, Takumi Fujimoto, Jo o Gabriel Oliveira, Isaac Galatzer-Levy, Jo o Gante, Simon Geisler, Soham Ghosal, Antonious M. Girgis, Tamara von Glehn, Alec Go, Alhaad Gokhale, Alex Grills, Yiming Gu, Mayank Gupta, Pramod Gupta, Guru Guruganesh, Raia Hadsell, Hamza Harkous, Jitendra Harlalka, Demis Hassabis, Anja Hauth, Joe Heyward, Arian Hosseini, Chih-Yang Hsia, I-Hung Hsu, Xiaopeng Huang, Yangsibo Huang, Kevin Hui, Adrian Hutter, Te I, Fotis Iliopoulos, Advait Jain, Ganesh Jawahar, Ziwei Ji, Qilin Jin, Melvin Johnson, Kandarp Joshi, Arun Kandoor, Wang-Cheng Kang, Koray Kavukcuoglu, Mehran Kazemi, Kathleen Kenealy, Amr

315 Khalifa, Phoebe Kirk, Ivan Korotkov, Suraj Kothawade, Vitaly Kovalev, Neel Kovelamudi, Adam
316 Kraft, Ravin Kumar, Vivek Kumar, Harish Kuppam, Justin Lannin, Chen-Yu Lee, Seungji Lee,
317 Dmitry Lepikhin, Alon Levkovitch, Dongdong Li, Qiujia Li, Valentin Liévin, Ethan Lin, Ziqian
318 Lin, Casper Liu, Tianlin Liu, Tianqi Liu, Xin Liu, Ivan Lobov, Mayank Lunayach, Min Ma, Gagan
319 Madan, Andrii Maksai, Eric Malmi, Michal Matuszak, Daniel McDuff, Gaurav Menghani, Maciej
320 Mikuła, Daniil Mirylenka, Karolis Misiunas, Vedant Misra, Andreea Mitran, Kareem Mohamed,
321 Maksim Mukha, Eric Noland, James O’Donnell, Brendan O’Donoghue, Kate Olszewska, Bernett
322 Orlando, Wanqiong Pan, Rina Panigrahy, Unnati Parekh, Nicolas Perez-Nieves, Chunjong Park,
323 Eric Paskie, Liqian Peng, Bryce Petrini, Slav Petrov, Jonas Pfeiffer, Bilal Piot, Martyna Plomecka,
324 Siim Poder, Octavio Ponce, Arijit Pramanik, David Racz, Anish Rajan, Michelle Ramanovich,
325 Anand Rao, Marvin Ritter, Vitor Rodrigues, Evan Rosen, Mikołaj Rybiński, Noveen Sachdeva,
326 Michaël E. Sander, Rohit Sathyanarayana, Sagar Savla, Samuel Schmidgall, Tal Schuster, George
327 Scrivener, Benoit Seguin, Andrew Sellergren, Aliaksei Severyn, Izhak Shafran, Dhruv Shah, Bobak
328 Shahriari, Yuan Shangquan, Ashish Shenoy, Pradeep Shenoy, Rakesh Shivanna, Pauline Sho,
329 Lucas Spangher, Wojciech Stokowiec, Tim Strother, Yao Su, Yinghao Sun, Mukund Sundararajan,
330 Andrea Tacchetti, Mor Hazan Taege, Pouya Tafti, Jean Tarbouriech, Chetan Tekur, Shantanu
331 Thakoor, Rahul Thapa, Madeleine Traverse, Lenart Treven, Tao Tu, Chien Te Tung, Çağlar Ünlü,
332 Petar Veličković, Malini Pooni Venkat, Sagar Gubbi Venkatesh, Vidya Venkiteswaran, Francesco
333 Visin, Alex Vitvitskyi, Kiran Vodrahalli, Weiyi Wang, Xin Wang, Tris Warkentin, Jan Wassenberg,
334 John Wieting, Cindy Wu, Lechao Xiao, Hao Xu, Yuhui Xu, Fuzhao Xue, Arun Yadav, Jun Yan,
335 Antoine Yang, Lin Yang, Ming-Hsuan Yang, Ziyu Ying, Jae Hyeon Yoo, Morteza Zadimoghaddam,
336 Sajjad Zafar, Fred Zhang, Jiageng Zhang, Jianyi Zhang, Xiaofan Zhang, Chao Zhao, David Zhou,
337 and Chen Zou. Gemma 4 technical report, 2026. URL <https://arxiv.org/abs/2607.02770>.

338 Jiahui Geng, Fengyu Cai, Yuxia Wang, Heinz Koepl, Preslav Nakov, and Iryna Gurevych. A
339 Survey of Confidence Estimation and Calibration in Large Language Models, 2024. URL <https://arxiv.org/abs/2311.08298>.

341 Sarthak Jain and Byron C. Wallace. Attention is not Explanation, 2019. URL <https://arxiv.org/abs/1902.10186>.

343 Zhewei Kang, Xuandong Zhao, and Dawn Song. Scalable Best-of-N Selection for Large Language
344 Models via Self-Certainty, 2025. URL <https://arxiv.org/abs/2502.18581>.

345 Hojin Kim and Jaehyung Kim. Reasoning or fluency? dissecting probabilistic confidence in best-of-n
346 selection, 2026. URL <https://arxiv.org/abs/2601.13735>.

347 Sicong Leng, Hang Zhang, Guanzheng Chen, Xin Li, Shijian Lu, Chunyan Miao, and Lidong Bing.
348 Mitigating Object Hallucinations in Large Vision-Language Models through Visual Contrastive De-
349 coding. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*
350 *(CVPR)*, pages 13872–13882, June 2024.

351 Yifan Li, Yifan Du, Kun Zhou, Jinpeng Wang, Wayne Xin Zhao, and Ji-Rong Wen. Evaluating Object
352 Hallucination in Large Vision-Language Models, 2023. URL <https://arxiv.org/abs/2305.10355>.

354 Hunter Lightman, Vineet Kosaraju, Yura Burda, Harri Edwards, Bowen Baker, Teddy Lee, Jan
355 Leike, John Schulman, Ilya Sutskever, and Karl Cobbe. Let’s Verify Step by Step, 2023. URL
356 <https://arxiv.org/abs/2305.20050>.

357 Pan Lu, Hritik Bansal, Tony Xia, Jiacheng Liu, Chunyuan Li, Hannaneh Hajishirzi, Hao Cheng,
358 Kai-Wei Chang, Michel Galley, and Jianfeng Gao. MathVista: Evaluating Mathematical Reasoning
359 of Foundation Models in Visual Contexts, 2024. URL <https://arxiv.org/abs/2310.02255>.

360 Fawaz Sammani, Tzoulis Chamiti, and Nikos Deligiannis. On Test-Time Scaling for Vision-Language
361 Models, 2026. URL <https://arxiv.org/abs/2606.28864>.

362 Charlie Snell, Jaehoon Lee, Kelvin Xu, and Aviral Kumar. Scaling LLM Test-Time Compute
363 Optimally can be More Effective than Scaling Model Parameters, 2024. URL <https://arxiv.org/abs/2408.03314>.

- Amir Taubenfeld, Tom Sheffer, Eran Ofek, Amir Feder, Ariel Goldstein, Zorik Gekhman, and Gal Yona. Confidence Improves Self-Consistency in LLMs. In Wanxiang Che, Joyce Nabende, Ekaterina Shutova, and Mohammad Taher Pilehvar, editors, *Findings of the Association for Computational Linguistics: ACL 2025*, pages 20090–20111, Vienna, Austria, July 2025. Association for Computational Linguistics. ISBN 979-8-89176-256-5. doi: 10.18653/v1/2025.findings-acl.1030. URL <https://aclanthology.org/2025.findings-acl.1030/>.
- Yijie Tong, Yifan Hou, Shaobo Cui, Antoine Bosselut, and Mrinmaya Sachan. Diversity Matters: Revisiting Test-Time Compute in Vision-Language Models, 2026. URL <https://arxiv.org/abs/2605.30713>.
- Xuezhi Wang, Jason Wei, Dale Schuurmans, Quoc Le, Ed Chi, Sharan Narang, Aakanksha Chowdhery, and Denny Zhou. Self-Consistency Improves Chain of Thought Reasoning in Language Models, 2023. URL <https://arxiv.org/abs/2203.11171>.
- Zengbin Wang, Feng Xiong, Liang Lin, Xuecai Hu, Yong Wang, Yanlin Wang, Man Zhang, and Xiangxiang Chu. Visually-Guided Policy Optimization for Multimodal Reasoning, 2026. URL <https://arxiv.org/abs/2604.09349>.
- Jason Wei, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, Brian Ichter, Fei Xia, Ed Chi, Quoc Le, and Denny Zhou. Chain-of-Thought Prompting Elicits Reasoning in Large Language Models, 2022. URL <https://arxiv.org/abs/2201.11903>.
- Sarah Wiegrefe and Yuval Pinter. Attention is not not explanation. In Kentaro Inui, Jing Jiang, Vincent Ng, and Xiaojun Wan, editors, *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, pages 11–20, Hong Kong, China, November 2019. Association for Computational Linguistics. doi: 10.18653/v1/D19-1002. URL <https://aclanthology.org/D19-1002/>.
- Zhangyue Yin, Qiushi Sun, Qipeng Guo, Zhiyuan Zeng, Xiaonan Li, Tianxiang Sun, Cheng Chang, Qinyuan Cheng, Ding Wang, Xiaofeng Mou, Xipeng Qiu, and XuanJing Huang. Aggregation of Reasoning: A Hierarchical Framework for Enhancing Answer Selection in Large Language Models, 2024. URL <https://arxiv.org/abs/2405.12939>.
- Xiang Yue, Tianyu Zheng, Yuansheng Ni, Yubo Wang, Kai Zhang, Shengbang Tong, Yuxuan Sun, Botao Yu, Ge Zhang, Huan Sun, Yu Su, Wenhui Chen, and Graham Neubig. MMMU-Pro: A More Robust Multi-discipline Multimodal Understanding Benchmark, 2025. URL <https://arxiv.org/abs/2409.02813>.
- Zhuosheng Zhang, Aston Zhang, Mu Li, Hai Zhao, George Karypis, and Alex Smola. Multimodal Chain-of-Thought Reasoning in Language Models, 2024. URL <https://arxiv.org/abs/2302.00923>.
- Jinguo Zhu, Weiyun Wang, Zhe Chen, Zhaoyang Liu, Shenglong Ye, Lixin Gu, Hao Tian, Yuchen Duan, Weijie Su, Jie Shao, Zhangwei Gao, Erfei Cui, Xuehui Wang, Yue Cao, Yangzhou Liu, Xingguang Wei, Hongjie Zhang, Haomin Wang, Weiye Xu, Hao Li, Jiahao Wang, Nianchen Deng, Songze Li, Yinan He, Tan Jiang, Jiapeng Luo, Yi Wang, Conghui He, Botian Shi, Xingcheng Zhang, Wenqi Shao, Junjun He, Yingting Xiong, Wenwen Qu, Peng Sun, Penglong Jiao, Han Lv, Lijun Wu, Kaipeng Zhang, Huipeng Deng, Jiaye Ge, Kai Chen, Limin Wang, Min Dou, Lewei Lu, Xizhou Zhu, Tong Lu, Dahua Lin, Yu Qiao, Jifeng Dai, and Wenhui Wang. InternVL3: Exploring Advanced Training and Test-Time Recipes for Open-Source Multimodal Models, 2025. URL <https://arxiv.org/abs/2504.10479>.

409 A Experimental and Implementation Details

410 **Prompt templates** Each prompt contains the image(s), question, and answer choices when applicable. Qwen2.5-VL and InternVL3 receive the following chain-of-thought prompt: “Think step by step inside <think> </think> tags, then give your final answer inside \boxed{ }.” Gemma4 is run with thinking mode disabled and receives the same instruction without explicit thinking tags: “Think step by step, then give your final answer inside \boxed{ }.” Inputs are rendered using each model’s native chat template.

Dataset information We evaluate the 1,000-example MathVista testmini split, the complete 1,730-example MMMU-Pro standard test set, and the complete 1,500-example MMStar validation set for all three models.

MathVista evaluates mathematical reasoning in visual contexts and contains both multiple-choice and free-form questions. It contains arithmetic, algebraic, geometric, logical, numerical, statistical, and scientific reasoning problems, with visual inputs including diagrams, charts, tables, function plots, scientific figures, natural images, and puzzle-style scenes [Lu et al., 2024].

MMMU-Pro targets college-level reasoning across a broad range of disciplines, including science, engineering, medicine, business, humanities, and the social sciences. Questions require combining domain knowledge with information from various visual inputs such as diagrams, plots, tables, maps, and scientific figures. The benchmark is constructed from MMMU by filtering questions that can be answered reliably without the image, making visual evidence more necessary for solving the task [Yue et al., 2025].

MMStar consists of visually dependent multiple-choice questions organized across six capability categories: coarse perception, fine-grained perception, instance reasoning, logical reasoning, science and technology, and mathematics. These are further divided into 18 fine-grained dimensions covering tasks such as object localization and counting, attribute recognition, cross-instance comparison, diagram reasoning, numerical calculation, and statistical analysis. Its examples are curated to reduce language-only shortcuts and ensure that visual information is necessary for answering the question [Chen et al., 2024].

Answer extraction and normalization We extract the last well-formed $\boxed{\cdot}$ expression in each response. Multiple-choice answers are canonicalized to uppercase option letters. Boxed option text is matched to its corresponding choice. When no valid boxed output is present, conservative final-sentence and final-line patterns are used. Free-form MathVista answers are compared using normalized numerical or string matching, including fraction handling and decimal comparison. Unextractable responses are assigned a shared <NO_ANSWER> label.

Tie handling Direct Best-of-N selection resolves score ties using the lowest original rollout index. For RWC, equal VisConf scores receive their average rank; if all scores are equal, every rollout receives weight 0.5. If answer-level votes are tied, we select the answer whose earliest supporting rollout has the lowest original index.

Compute Resources Experiments with Qwen2.5-VL-3B and Gemma4-E2B used one NVIDIA GeForce RTX 4090 GPU with 24 GB of memory; InternVL3-8B experiments used one NVIDIA A100 GPU with 80 GB of memory.

Table 3: LookBack Best-of-N and Rank-Weighted Consensus (RWC) performance under two values of α , with $\lambda = 1.25$.

α	λ	Aggregation	Qwen2.5-VL-3B			InternVL3-8B			Gemma4-E2B			Avg.
			$N = 8$	$N = 16$	$N = 32$	$N = 8$	$N = 16$	$N = 32$	$N = 8$	$N = 16$	$N = 32$	
0.25	1.25	BoN	49.88	50.17	50.45	56.31	55.00	52.99	48.92	49.38	49.95	51.45
0.25	1.25	RWC	53.40	54.61	55.51	60.78	61.68	62.07	50.16	50.70	50.84	55.53
0.50	1.25	BoN	50.09	50.34	50.57	58.22	58.11	57.84	49.16	49.74	50.27	52.70
0.50	1.25	RWC	53.68	54.93	55.72	61.27	62.16	62.65	50.30	50.86	51.19	55.86

LookBack Implementation For LookBack, let $\ell_{i,t}$ be the sampled-token log-probability and $A_{i,t}$ the all-layer, all-head image-attention mass. Its response score is

$$L_i = \frac{\sum_t A_{i,t}^\lambda (\ell_{i,t} + \alpha \log A_{i,t})}{\sum_t A_{i,t}^\lambda}.$$

The work suggests a choice of $\alpha = 0.25$ and $\lambda = 1.25$ as a global setting. However, we observe that $\alpha = 0.5$ and $\lambda = 1.25$, which were the model-specific settings for Qwen2.5-VL-7B, achieve greater performance (Table 3). Thus, we fix $\alpha = 0.5$ and $\lambda = 1.25$ across all models, datasets, and sampling budgets in our baseline results. Standard LookBack selects the highest-scoring rollout. LookBack + RWC applies our rank-weighted answer aggregation to the same L_i scores.

B Full Results and Statistical Analysis

B.1 Per-Model, Dataset, and N Results

Table 4: Full accuracy results (%) on MathVista, MMMU-Pro, and MMStar for $N \in \{8, 16, 32\}$. Avg. is the unweighted mean across the nine dataset-budget combinations within each model.

Method	MathVista			MMMU-Pro			MMStar			Avg.
	$N = 8$	$N = 16$	$N = 32$	$N = 8$	$N = 16$	$N = 32$	$N = 8$	$N = 16$	$N = 32$	
Qwen2.5-VL-3B										
Oracle	82.76	87.01	89.40	80.31	88.10	92.95	85.50	91.57	95.67	88.14
FirstAns	54.01	54.01	54.01	41.20	41.20	41.20	50.31	50.31	50.31	48.51
Self-Consistency	61.04	62.48	63.20	45.30	46.38	46.88	54.77	55.38	55.13	54.51
Self-Certainty	55.65	55.70	55.40	42.97	43.25	43.76	50.72	50.89	50.67	49.89
+ Borda ($p = 0.3$)	61.35	62.78	63.60	45.52	46.55	46.76	54.63	55.27	55.07	54.61
+ Borda ($p = 0.7$)	61.25	62.79	63.90	45.60	46.64	47.28	54.48	55.08	54.93	54.66
+ Borda ($p = 1.2$)	60.88	62.61	63.90	45.55	46.58	47.23	54.15	54.96	54.87	54.53
+ Borda ($p = 2.0$)	60.07	62.26	63.90	45.28	46.42	47.23	53.68	54.72	55.07	54.29
LookBack	55.56	55.64	55.50	42.98	43.22	43.41	51.74	52.17	52.80	50.33
+ RWC	60.66	62.44	63.70	45.75	46.90	47.34	54.64	55.44	56.13	54.78
VisConf	56.63	56.82	57.10	43.57	43.44	43.06	52.93	53.21	53.33	51.12
+ RWC	61.08	62.97	64.50	46.12	47.24	48.38	55.06	55.75	56.00	55.23
InternVL3-8B										
Oracle	80.46	83.92	86.50	80.69	86.67	90.64	84.77	88.81	91.80	86.03
FirstAns	57.57	57.57	57.57	49.43	49.43	49.43	62.65	62.65	62.65	56.55
Self-Consistency	63.62	64.66	64.60	53.28	54.02	54.34	65.97	66.44	66.47	61.49
Self-Certainty	60.75	60.79	60.50	51.68	52.15	52.37	64.05	64.16	64.33	58.98
+ Borda ($p = 0.3$)	64.05	65.12	65.70	53.55	54.24	54.45	66.03	66.55	66.60	61.81
+ Borda ($p = 0.7$)	64.15	65.25	66.10	53.66	54.38	54.62	66.06	66.57	66.73	61.95
+ Borda ($p = 1.2$)	63.99	65.19	66.40	53.62	54.43	55.14	66.03	66.49	66.80	62.01
+ Borda ($p = 2.0$)	63.58	64.91	66.20	53.41	54.33	55.20	65.83	66.45	66.60	61.84
LookBack	59.46	59.34	59.80	50.85	50.82	50.17	64.34	64.17	63.53	58.05
+ RWC	63.71	65.00	65.90	53.72	54.60	54.97	66.38	66.88	67.07	62.02
VisConf	61.04	61.26	61.40	52.25	52.90	53.41	65.10	65.27	66.00	59.85
+ RWC	64.29	65.49	67.00	54.13	54.96	55.55	66.65	67.14	67.27	62.50
Gemma4-E2B										
Oracle	74.64	80.08	83.50	72.81	79.87	84.97	79.38	84.96	88.73	80.99
FirstAns	47.63	47.63	47.63	41.90	41.90	41.90	52.49	52.49	52.49	47.34
Self-Consistency	50.14	50.56	50.90	44.29	44.59	44.68	55.10	55.55	55.47	50.14
Self-Certainty	49.21	49.43	50.60	43.77	43.87	43.76	55.30	55.55	55.07	49.62
+ Borda ($p = 0.3$)	50.34	50.76	50.90	44.49	44.83	45.38	55.54	55.90	55.87	50.44
+ Borda ($p = 0.7$)	50.44	50.89	50.80	44.62	45.03	45.49	55.75	56.18	56.13	50.59
+ Borda ($p = 1.2$)	50.49	50.90	50.60	44.69	45.16	45.32	55.87	56.38	56.80	50.69
+ Borda ($p = 2.0$)	50.44	50.87	50.80	44.71	45.23	45.78	55.94	56.54	56.87	50.80
LookBack	48.98	49.40	49.80	44.00	44.97	46.01	54.48	54.85	55.00	49.72
+ RWC	50.35	50.76	51.10	45.00	45.66	46.07	55.54	56.15	56.40	50.78
VisConf	49.95	50.54	50.80	44.87	45.25	45.20	55.39	55.68	55.73	50.38
+ RWC	50.93	51.48	51.70	45.64	46.22	46.53	56.09	56.63	56.80	51.34

B.2 Paired Statistical Analysis

We estimate uncertainty using a paired question-cluster bootstrap (Table 5). For $N \in \{8, 16\}$, accuracy is first averaged over the 100 shared candidate subsets for each question; $N = 32$ uses the complete rollout pool. We resample the 4,230 questions with replacement within each benchmark,

Table 5: Paired accuracy improvements of VisConf + RWC over each comparator in percentage points with 95% question-cluster bootstrap confidence intervals. Overall averages across sampling budgets, models, and benchmarks.

Comparator	$N = 8$	$N = 16$	$N = 32$	Overall
LookBack + RWC	+0.47 [0.35, 0.59]	+0.45 [0.31, 0.61]	+0.56 [0.24, 0.87]	+0.49 [0.32, 0.66]
Self-Consistency	+0.72 [0.56, 0.86]	+0.87 [0.68, 1.07]	+1.34 [0.97, 1.72]	+0.98 [0.76, 1.18]
Self-Cert. + Borda ($p = 1.2$)	+0.52 [0.40, 0.66]	+0.58 [0.40, 0.74]	+0.74 [0.39, 1.09]	+0.61 [0.43, 0.80]
VisConf (top-1)	+2.03 [1.79, 2.28]	+2.61 [2.21, 3.02]	+3.08 [2.42, 3.77]	+2.57 [2.15, 3.00]

keeping evaluations of the same question paired across models. Method differences are computed before aggregation and macro-averaged across the nine model-benchmark combinations. The overall estimate additionally averages across $N \in \{8, 16, 32\}$. We report percentile 95% confidence intervals from 2,000 bootstrap replicates.

C Computational Overhead of Attention Extraction

We measure the overhead of replacing FlashAttention-2 with eager attention and explicit attention extraction using a controlled replay experiment on Qwen2.5-VL-3B. We replay the exact stored output-token sequences for nested candidate pools of $N \in \{8, 16, 32\}$ on 120 MMMU-Pro prompts, with three repetitions per condition, while holding prompts, images, generated tokens, batching, and KV caching fixed. Figure 5 reports the paired wall-clock runtime ratio of eager attention extraction to FlashAttention-2, computed per prompt after taking the median over repetitions. Median runtime is comparable at $N = 8$ and $N = 16$, with ratios of $0.99\times$ and $1.01\times$, respectively, while the overhead becomes more pronounced at $N = 32$, where eager attention extraction requires $1.18\times$ the runtime. Because the benchmark excludes scoring and answer aggregation, these measurements represent the inference-time cost of obtaining the attention signal itself.




Figure 5: Paired wall-clock runtime ratio of eager attention extraction to FlashAttention-2 on Qwen2.5-VL-3B across 120 MMMU-Pro prompts. Points show median ratios and error bars show the interquartile range.

D Analysis of the VisConf Score


D.1 Joint Enrichment of Confidence and Visual Engagement

Figure 6 examines each signal while stratifying by the other. Rollout accuracy generally increases with both confidence and visual engagement, while weak values of either signal limit accuracy. This complementary enrichment motivates combining the two signals in VisConf.

D.2 Token- and Trajectory-Level Fusion

We analyze the raw confidence signal $k_{i,t}$, before applying the candidate-relative calibration used by VisConf, to isolate the effect of where confidence and visual engagement are fused. For rollout i with T_i generated tokens, define

$$V_i = \frac{1}{T_i} \sum_{t=1}^{T_i} v_{i,t}, \quad K_i = \frac{1}{T_i} \sum_{t=1}^{T_i} k_{i,t}. \quad (1)$$



(a) Accuracy across confidence percentiles, (b) Accuracy across visual engagement percentiles, stratified by confidence tercile.

Figure 6: Joint enrichment of rollout correctness by raw confidence and visual engagement. The dashed line denotes overall rollout accuracy, and shaded regions show 95% confidence intervals.

Token-level and trajectory-level fusion are respectively

$$Q_i^{\text{tok}} = \frac{1}{T_i} \sum_{t=1}^{T_i} v_{i,t} k_{i,t}, \quad Q_i^{\text{traj}} = V_i K_i. \quad (2)$$

Covariance decomposition By the definition of within-rollout covariance,

$$Q_i^{\text{tok}} = V_i K_i + \text{Cov}_t(v_{i,t}, k_{i,t}). \quad (3)$$

Token-level fusion therefore augments the product of the two trajectory means with a term determined by their temporal co-occurrence. Trajectory-level fusion retains the marginal strength of each signal without requiring them to be large at the same decoding positions.

Permutation identity Let π_i be a uniformly random permutation of the T_i visual engagement positions. Then

$$\mathbb{E}_{\pi_i} \left[\frac{1}{T_i} \sum_{t=1}^{T_i} v_{i,\pi_i(t)} k_{i,t} \right] = V_i K_i. \quad (4)$$

For every position t , the permuted value $v_{i,\pi_i(t)}$ is uniformly distributed over the visual engagement sequence and therefore has expectation V_i . Consequently,

$$\frac{1}{T_i} \sum_t \mathbb{E}_{\pi_i}[v_{i,\pi_i(t)}] k_{i,t} = \frac{1}{T_i} \sum_t V_i k_{i,t} = K_i V_i. \quad (5)$$

Thus, trajectory-level fusion is the expected token product after temporal correspondence between confidence and visual engagement has been removed. The same identities hold after replacing raw confidence $k_{i,t}$ with calibrated confidence $c_{i,t}$; VisConf uses the resulting trajectory-level product $V_i C_i$.

Within-question discrimination We evaluate these formulations on the 3,558 mixed-correctness questions from Qwen2.5-VL-3B across MathVista, MMMU-Pro, and MMStar at $N = 32$. A mixed-correctness question contains at least one correct and one incorrect rollout.

For question q , let $\mathcal{C}_q$ and $\mathcal{I}_q$ denote its sets of correct and incorrect rollouts. For rollout score s , we compute

$$\text{AUC}_q(s) = \frac{1}{|\mathcal{C}_q| |\mathcal{I}_q|} \sum_{i \in \mathcal{C}_q} \sum_{j \in \mathcal{I}_q} \mathbf{1}[s_i > s_j]. \quad (6)$$

We macro-average this quantity across questions so that questions with more correct–incorrect rollout pairs do not receive greater weight. Confidence intervals are obtained by bootstrap resampling questions. For permuted token fusion, we average over 100 independently sampled within-rollout permutations.




Figure 7: Post-hoc sensitivity of VisConf to the confidence-vision weighting parameter α . Performance is strongest near the balanced setting $\alpha = 0.5$.




Figure 8: Answer correctness by same-answer support and VisConf rank. Accuracy increases along both dimensions.

Table 6: Within-question discrimination of rollout-quality signals on Qwen2.5-VL-3B, pooled across MathVista, MMMU-Pro, and MMStar. Fusion-location controls use raw confidence k , whereas VisConf uses candidate-calibrated confidence c .

Signal	Score	AUC	95% CI
Confidence	K_i	0.5279	[0.5214, 0.5344]
Visual engagement	V_i	0.5307	[0.5243, 0.5374]
Aligned token fusion	$\frac{1}{T_i} \sum_t k_{i,t} v_{i,t}$	0.5373	[0.5304, 0.5440]
Permuted token fusion	$\frac{1}{T_i} \sum_t k_{i,t} v_{i,\pi_i(t)}$	0.5399	[0.5333, 0.5468]
Trajectory fusion	$K_i V_i$	0.5399	[0.5331, 0.5462]
VisConf	$C_i V_i$	0.5503	[0.5434, 0.5569]

498

499 Randomizing token correspondence produces essentially the same discrimination as trajectory-level
500 fusion, whereas preserving the original token alignment is slightly weaker. Tokenwise correspondence
501 therefore provides no additional discrimination in our analysis.

502 The covariance term in Eq. 3 is negatively oriented with respect to correctness. Its within-question
503 AUC is 0.4634, with a 95% confidence interval of [0.4570, 0.4697]. The question-macro mean
504 covariance is -0.02659 for correct rollouts and -0.02507 for incorrect rollouts, giving a correct-
505 minus-incorrect difference of -0.00152 , with a 95% confidence interval of $[-0.00193, -0.00110]$.

506 These observations are consistent with the lack of improvement from adding the observed covariance
507 term. They do not imply that correct reasoning requires confidence and visual processing to occur at
508 separate stages. Rather, they show that greater tokenwise co-occurrence may not be a useful quality
509 signal.

510 D.3 Sensitivity to Confidence-Vision Weighting

511 To test the sensitivity of VisConf to the relative contribution of confidence and visual engagement,
512 we generalize its multiplicative score as

$$Z_{i,\alpha} = V_i^\alpha C_i^{1-\alpha}, \quad \alpha \in [0, 1].$$

513 This weighted geometric form preserves the requirement that both signals contribute while allowing
514 their relative importance to vary. The endpoints $\alpha = 0$ and $\alpha = 1$ recover confidence-only and vision-
515 only scoring, respectively. At $\alpha = 0.5$, $Z_{i,0.5} = \sqrt{V_i C_i}$, which is rank-equivalent to the VisConf
516 product $V_i C_i$. The post-hoc sweep in Figure 7 shows that performance is strongest under broadly
517 balanced weighting and varies smoothly around the symmetric choice, indicating that VisConf is not
518 narrowly dependent on $\alpha = 0.5$.

E Analysis of Rank-Weighted Consensus

E.1 Answer Support and Rollout Quality

RWC combines answer frequency with rollout quality. Figure 8 shows that, at fixed same-answer support, higher VisConf-ranked rollouts are more accurate. Likewise, at fixed rank, greater answer support is associated with higher accuracy. This supports combining both signals in RWC.

E.2 Rescues and Regressions by Rollout Correctness Frequency

Table 7: Rescues and regressions of VisConf + RWC relative to Self-Consistency on InternVL3-8B at $N = 32$, pooled across MathVista, MMMU-Pro, and MMStar.

Correct rollouts	Questions	Rescues	Regressions	Net
1–4	521	6	1	+5
5–8	355	18	1	+17
9–12	340	52	17	+35
13–15	197	34	10	+24
1–15	1,413	110	29	+81
16–32	2,397	0	24	−24

To localize the paired changes underlying the gains in Figure 4, we partition questions by the number of correct rollouts in the $N = 32$ pool. A *rescue* occurs when Self-Consistency is wrong but VisConf + RWC is correct; a *regression* is the reverse. Table 7 reports these transitions on the same InternVL3-8B evaluation used in Figure 4.

E.3 Robustness Properties of Rank-Weighted Consensus

RWC depends only on the ordering of VisConf scores. Consequently, any strictly increasing transformation of the scores leaves the rollout weights and selected answer unchanged. This makes aggregation insensitive to the absolute scale of VisConf.

For $w_i = (N - r_i)/(N - 1)$, the rank weights also satisfy

$$\sum_{i=1}^N w_i = \frac{N}{2}, \quad \frac{w_i}{\sum_j w_j} \leq \frac{2}{N}.$$

Thus, no individual rollout contributes more than $2/N$ of the normalized RWC mass. An isolated extreme score therefore cannot dominate the decision through its magnitude; an answer must instead accumulate support across high-ranked trajectories.

E.4 Borda and RWC Aggregation Rules

Table 8: Accuracy (%) on Qwen2.5-VL-3B, obtained by applying RWC or parameterized Borda aggregation to the same VisConf rollout scores. Avg is the unweighted mean across MathVista, MMMU-Pro, and MMStar for $N \in \{8, 16, 32\}$.

Aggregation	MathVista			MMMU-Pro			MMStar			Avg.
	$N = 8$	$N = 16$	$N = 32$	$N = 8$	$N = 16$	$N = 32$	$N = 8$	$N = 16$	$N = 32$	
RWC	61.08	62.97	64.50	46.12	47.24	48.38	55.06	55.75	56.00	55.23
Borda ($p = 0.3$)	61.46	63.02	64.10	45.79	46.80	47.28	55.00	55.59	55.73	54.97
Borda ($p = 0.7$)	61.46	63.11	64.40	46.04	47.09	48.15	55.10	55.69	55.87	55.21
Borda ($p = 1.2$)	61.14	62.91	64.10	46.10	47.20	48.38	55.07	55.74	56.07	55.19
Borda ($p = 2.0$)	60.41	62.43	63.60	45.92	47.12	48.32	54.85	55.73	56.00	54.93

538 Parameterized Borda aggregation assigns

$$b_p(r_i) = (N - r_i + 1)^p.$$

539 For $p = 1$, this is an affine transformation of the RWC rollout weight:

$$b_1(r_i) = (N - 1)w_i + 1.$$

540 Consequently, its answer-group score satisfies

$$S_{p=1}(a) = (N - 1)S_{\text{RWC}}(a) + n_a,$$

541 Therefore, linear Borda adds an answer-frequency term to the rank-weighted quality evidence. This
 542 distinction is relevant to the failure mode studied in Section 7.3 when correct trajectories are high-
 543 quality but remain a minority: the additive n_a term in linear Borda explicitly favors the more frequent
 544 answer, partially restoring the majority bias that RWC is intended to overcome. More generally,
 545 $p < 1$ further compresses rank differences and approaches unweighted consensus as $p \rightarrow 0$, whereas
 546 $p > 1$ increasingly concentrates influence on the highest-ranked trajectories. RWC instead uses a
 547 fixed linear rank schedule with no additive frequency bonus.

548 Across the range of evaluated p values, Borda aggregation is broadly comparable to RWC, reinforcing
 549 that rank-weighted answer aggregation is generally effective for combining rollout quality with
 550 consensus. RWC offers a simpler alternative by fixing the weighting schedule and removing the
 551 need to tune the exponent p , while retaining comparable performance across models, datasets, and
 552 sampling budgets.

553 F Qualitative Examples

554 We present five examples at $N = 32$ to illustrate distinct qualitative outcomes: two corrections by
 555 VisConf + RWC, one regression, one case in which all methods succeed, and one case in which all
 556 methods fail. The examples span two models and all three benchmarks. For each example, *support*
 557 denotes either the number of votes for the selected answer under Self-Consistency or the sum of
 558 normalized within-question rank weights under RWC. LookBack top-1 selects a single rollout and
 559 therefore has no answer-group support value.

560 F.1 Corrections by VisConf + RWC

561

(a) Visual proportion estimation – Qwen2.5-VL-3B, MMStar sample 359



Question. Roughly how much of the image is occupied by the two people on the bench? The relevant options are A: more than 50% and B: less than 30%.

Gold answer: B (less than 30%)

Candidate pool. A receives 16 rollouts, B receives 15, and D receives one.

562

Method	Prediction	Outcome	Support
Self-Consistency	A	incorrect	16 votes
LookBack top-1	A	incorrect	–
LookBack + RWC	A	incorrect	8.48
VisConf + RWC	B	correct	9.06

Interpretation. The unweighted vote is wrong by one rollout. Correct trajectories explicitly compare the people with the much larger water and sky regions, whereas many A trajectories treat the people as the image’s semantic focus and infer that they must occupy most of its area. VisConf + RWC gives the minority B group enough collective weight to reverse this narrow incorrect plurality.

563

(b) Tacheometric calculation – InternVL3-8B, MMMU-Pro sample
test_Architecture_and_Engineering_492

Instrument station	Staff station	Vertical angle	Staff readings (m)	Remarks (m)
	B.M.	$-6^{\circ}0'$	1.100, 1.153, 2.060	R.L. of B.M.
A	B	$+8^{\circ}0'$	0.982, 1.105, 1.188	= 976.00 m

Question. Using the staff readings in the figure, determine the horizontal distance from A to B and the reduced level of B.

Gold answer: D: $AB = 94.14$ m and R.L. = 988.81 m

Candidate pool. C receives 14 rollouts, D receives 13, A and no-answer receive two each, and B receives one.

564

Method	Prediction	Outcome	Support
Self-Consistency	C	incorrect	14 votes
LookBack top-1	C	incorrect	–
LookBack + RWC	C	incorrect	7.03
VisConf + RWC	D	correct	7.61

Interpretation. The candidate pool is fragmented and the correct answer trails the leading incorrect answer by one rollout. The incorrect C trajectories typically reach a nearby horizontal-distance value but propagate an error through the trigonometric or reduced-level calculation. VisConf + RWC shifts the decision to the D group, whose trajectories more consistently retain the required quantities through the complete calculation.

565 **F.2 A Regression**

566

(c) Relational counting – InternVL3-8B, MathVista sample 984



Question. Is the number of gray cars less than the number of small metallic minivans? The options are A: yes and B: no.

Gold answer: B (no)

Candidate pool. B receives 17 rollouts and A receives 15.

567


Method	Prediction	Outcome	Support
Self-Consistency	B	correct	17 votes
LookBack top-1	B	correct	–
LookBack + RWC	B	correct	8.58
VisConf + RWC	A	incorrect	9.58

Interpretation. Here, the calibrated score is misleading. The candidate pool has a narrow but correct majority for B, and both LookBack variants preserve that decision. VisConf instead ranks enough of the incorrect A trajectories highly to reverse the result. This example illustrates that confidence and visual engagement remain imperfect quality signals: strong engagement with an image does not guarantee that the depicted objects were counted or categorized correctly.

568 **F.3 Agreement and Shared Failure**

569

(d) Mood recognition – Qwen2.5-VL-3B, MMStar sample 13



Question. Which mood does the image convey? The relevant options are B: cozy and C: happy.

Gold answer: B (cozy)

Candidate pool. B receives 16 rollouts, C receives 15, and one rollout produces no answer.

570

Method	Prediction	Outcome	Support
Self-Consistency	B	correct	16 votes
LookBack top-1	B	correct	–
LookBack + RWC	B	correct	9.61
VisConf + RWC	B	correct	8.23

Interpretation. This is a closely contested but stable agreement case. Both leading groups recognize the warm sunset and the companionship between the adult and child; they differ mainly in whether those cues imply “cozy” or “happy.” Despite the one-vote margin, both rollout-quality scores and both aggregation rules retain the correct B decision.

571

(e) Dominant-color recognition – Qwen2.5-VL-3B, MMStar sample 276



Question. What is the dominant color in the image? The relevant options are B: brown and C: black.

Gold answer: B (brown)

Candidate pool. C receives 17 rollouts and B receives 15. Thus, the pool contains a substantial minority of correct trajectories.

572

Method	Prediction	Outcome	Support
Self-Consistency	C	incorrect	17 votes
LookBack top-1	C	incorrect	–
LookBack + RWC	C	incorrect	9.94
VisConf + RWC	C	incorrect	8.87

Interpretation. This is a shared failure despite the presence of substantial correct evidence. The C trajectories emphasize individually black objects, such as clothing, instruments, and acoustic panels, while the B trajectories attend to the image-wide brown or sepia cast. Both quality scores reinforce the narrow incorrect majority rather than recovering the correct minority.

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [Yes]

Justification: The abstract and introduction summarize the paper’s contributions and empirical scope.

Guidelines:

- The answer [N/A] means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A [No] or [N/A] answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [Yes]

Justification: The Discussion explicitly addresses limitations.

Guidelines:

- The answer [N/A] means that the paper has no limitation while the answer [No] means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate “Limitations” section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren’t acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [N/A]

Justification: The paper focuses on empirical analysis and does not present formal theoretical results requiring theorem proofs.

Guidelines:

- The answer [N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes]

Justification: The Experimental Setup and Appendix describe the evaluated models and datasets, decoding configuration, sampling protocol, prompts, answer extraction and normalization, tie handling, baseline implementations, and the proposed scoring and aggregation procedures.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
 - (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [Yes]

Justification: The experiments use publicly available models and benchmarks, and an anonymized implementation is included in the current submission.

Guidelines:

- The answer [N/A] means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

Answer: [Yes]

Justification: The Experimental Setup and Appendix specify the required benchmark splits, model configurations, prompts, decoding hyperparameters, candidate budgets, subset-sampling procedure, baseline configurations, and evaluation protocol.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [Yes]

Justification: Appendix reports paired 95% confidence intervals using a question-cluster bootstrap with 2,000 replicates and describes the resampling and aggregation procedure.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The authors should answer [Yes] if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.

- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [Yes]

Justification: Appendix covers details about GPU model and memory used.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [Yes]

Justification: The research uses existing pretrained models and established public evaluation benchmarks and does not involve collection of personal data or experiments with human subjects. The research was conducted in accordance with the NeurIPS Code of Ethics.

Guidelines:

- The answer [N/A] means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer [No], they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [Yes]

Justification: Potential positive and negative impacts are discussed throughout the paper.

Guidelines:

- The answer [N/A] means that there is no societal impact of the work performed.
- If the authors answer [N/A] or [No], they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate Deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

Answer: [N/A]

Justification: The work does not introduce or release a new model or dataset which could require safeguards.

Guidelines:

- The answer [N/A] means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [Yes]

Justification: The code, data, and models used in the paper are properly cited and comply with the relevant licenses.

Guidelines:

- The answer [N/A] means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.

- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, paperswithcode.com/datasets has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [N/A]

Justification: The paper does not introduce a new dataset or model.

Guidelines:

- The answer [N/A] means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [N/A]

Justification: The research does not involve crowdsourcing experiments or research with human subjects.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. Institutional review board (IRB) approvals or equivalent for research with human subjects

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: [N/A]

Justification: The research does not involve crowdsourcing or human-subject experiments and therefore does not require IRB approval for the experiments reported in this work.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.

- 886 • Depending on the country in which research is conducted, IRB approval (or equivalent)
887 may be required for any human subjects research. If you obtained IRB approval, you
888 should clearly state this in the paper.
- 889 • We recognize that the procedures for this may vary significantly between institutions
890 and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the
891 guidelines for their institution.
- 892 • For initial submissions, do not include any information that would break anonymity (if
893 applicable), such as the institution conducting the review.

894 16. Declaration of LLM usage

895 Question: Does the paper describe the usage of LLMs if it is an important, original, or
896 non-standard component of the core methods in this research? Note that if the LLM is used
897 only for writing, editing, or formatting purposes and does *not* impact the core methodology,
898 scientific rigor, or originality of the research, declaration is not required.

899 Answer: [N/A]

900 Justification: The core method development in this research does not involve LLMs.

901 Guidelines:

- 902 • The answer [N/A] means that the core method development in this research does not
903 involve LLMs as any important, original, or non-standard components.
- 904 • Please refer to our LLM policy in the NeurIPS handbook for what should or should not
905 be described.